

06_project / AI / agentic web / August 2026

Second Audience

What an AI agent makes of your website, and what that costs

Your website has two kinds of visitor now. One of them has eyes. The other never runs your JavaScript, reads only what your server sends, and then answers questions about you for people who never open your page.

We measure what that second visitor can actually do, and since August 2026 we also measure what four to five agents make of it in practice. That second half is new, it operates under strict rules, and those rules are on page 4.

What we mean by "agent"

A system such as ChatGPT, Claude, Gemini, Perplexity or the AI answers in Google Search: software that reads web pages on someone's behalf and then answers in its own words. Wherever we write "agent", that is what is meant. We do not use the word "assistant".

The delta

Every page fetched twice, with and without JavaScript. What disappears is what the non-rendering crawlers cannot see.

The evidence

Every finding is a piece of text on a page, re-fetchable in the client's presence. Agent answers sit beside it as illustration, never as the finding.

The honesty

We name the layers that are theatre, write our prediction down before the test, and publish the places where we were wrong.

// the problem

You use it every day. Someone is already answering about you.

For many people this is now the first step rather than the last: asking a question instead of searching. The answer arrives in one paragraph, confidently, and is usually believed. What almost nobody does is turn the question around. Someone is asking it about you, today, and that answer comes from somewhere. This scan measures where.

This is not a technology that is coming. It is a habit you already have. The only new part is that you are on the other side of the question.

The blind spot in the technology

Most AI crawlers do not run JavaScript

Vercel measured this across its network: GPTBot, ClaudeBot, PerplexityBot, Meta-ExternalAgent, Bytespider and CCBot all fetch script files and run none of them. Content that only appears once the browser has drawn it does not score badly for those visitors. It does not exist.

The blind spot in the numbers

Your analytics structurally cannot show you this

Google Analytics loads through a script, so a crawler that never runs scripts cannot be reported. And the one category that is visible, agents driving a real Chrome, is precisely the category a server log cannot tell apart from a human. Two blind spots, and most audits cover one.

The blind spot in blocking

robots.txt is about fetching, not about answering

One provider's crawlers do not all do the same thing: GPTBot fetches for training, OAI-SearchBot builds the index ChatGPT quotes from, and ChatGPT-User fetches when a person asks for a page. Anyone who "blocks the AI bots" almost always blocks one of the three and does not know which. We meet this more often as an accident than as a decision.

The blind spot we had ourselves

What an agent then makes of it, we now measure too

Until August 2026 we stopped at what the server sends. That was defensible and it was too narrow: it left the most expensive question unanswered, namely whether the gap we point at actually gets filled with something wrong. We measure that now, under the rules on page 4.

// what it is

A scan that measures first, then the changes that travel

Five layers, and every finding comes with how we know.

Layer 1 · Rendering	The page fetched twice and compared. The difference is what the non-rendering crawlers cannot reach.
Layer 2 · Readability	Where each business fact actually lives: in text, in markup, in an image, or nowhere. And whether the same fact appears a second time elsewhere on the domain, worded differently.
Layer 3 · Permission	The intent of robots.txt against its effect, and the blocks nobody set on purpose.
Layer 4 · Usability	Whether an agent completes the task a visitor arrived to do, and which part of the answer came from the client's own pages.
Layer 5 · Proof	Thirty to ninety days of server log, classified, with claimed user agents verified against each provider's published address ranges. We advise against this layer more often than for it: it needs access to systems most clients do not run themselves, and at modest volumes it returns too little.

// the verdict

One question per task, not a score out of a hundred

Can an agent complete this today, what exactly blocks it, and what is the smallest change that fixes it. A verdict per task can be checked. A score invites gaming precisely the thing it claims to measure. The verdicts are named and the set grows with what we meet: answered by you · answered, but not by you · answered without the conditions · answered with the wrong figure · answered from the wrong page · answered with placeholder text · answered with an invented example · not reached · not answered.

When an AI agent arrives on your site instead of a person: what can it read, what can it do, and how would you know either? The scan is built to answer that from evidence rather than from a checklist.

// the method

Seven steps, in this order, and the agents come last

The order is the product. A scan that starts at step four finds technical imperfections and no money. Every strong finding we delivered this year came out of step three.

1 Identify the revenue object

Not the homepage. The thing that is sold: a night, a placement, a mandate, a subscription. Everything else is context. At one museum it turned out not to be admission but the €299 tower suite.

2 Extract the facts a buyer needs before committing

Price, availability, proof, eligibility, how to buy. For each one we record the literal sentence and the URL. That becomes the yardstick everything afterwards is scored against.

3 Hunt for the second answer

Does the same fact appear somewhere else on the domain, worded differently? In an FAQ, a stale translation, an orphan page, a test page, a press release from four years ago, or a dead address still in the search results. This is the engine of the scan.

across seven scans this step produced the headline finding in five of them

4 Check the form, not only the content

Is the fact there as text, as an image, or as prose with an escape word in it: "from", "typically", "average pricing"? A logo with no text beside it does not exist for a machine. A price that begins with "from" is an invitation to fill in the rest.

5 Map the retrieval layer

What do third parties say, literally, about the same fact? Guides, industry portals, company registers, LinkedIn. That is the material the gap gets filled with as soon as the client supplies nothing, and it can be quoted verbatim.

6 Test the finding against what has already been decided

Before a single agent is asked anything, search the published rulings, regulator guidance and enforcement decisions covering the class of fact the finding belongs to. Wording that is mandatory in the market it was written for must never be written up as an error.

on one scan this step removed a finding we had already drafted, and produced the strongest support in the document

7 Only then the agents, with the prediction written down first

Four to five agents, the same questions, each in a fresh session with no prior context, always followed by: at exactly which URL does that appear? We write our prediction down before the first question, so a hit is a hit and a miss does not quietly disappear.

One automatable rule that came out of step four and that we now run as standard: find brand images with no adjacent text and no meaningful alt attribute, and flag them as invisible proof. On one scan that single check predicted, in advance, that not one of sixteen agent answers would name a client. It was right.

```
// the rules at step seven
```

What we used to refuse to do, we now do, and under conditions

Until August 2026 we asked no agent anything about a client, because an AI answer cannot be reproduced. That argument still holds. The conclusion drawn from it did not: something that cannot be reproduced can still be recorded, and it is exactly what the client wants to see. Five rules keep it honest.

Rule 1

The agent answer is never the finding

The finding is always a piece of text on a page, still there tomorrow, re-fetchable in the client's presence. The agent answer sits beside it as an illustration of the consequence. That sentence appears verbatim in every report.

Rule 2

Always the follow-up question about the source

At exactly which URL does that appear? That one question turns the test from an impression into evidence. It produces two kinds of result, both useful: an agent naming a URL that does not contain the fact, and an agent correcting itself.

Rule 3

Fresh sessions, and the prediction written first

Every question in a clean window with no history, and the expectation on paper before the first question is asked. Without that, any outcome can be explained after the fact.

Rule 4

Our own model goes in unedited

Claude sits in the same table as the rest, with the same marks. On the most recent scan it earned credit for being the only one to refuse to give advice, and a red mark for contradicting itself on the same day about the client's city. Both are in there.

Rule 5

Our own misses go in the report

On one scan four of our five predictions were wrong. On another, two of five. That is in the document, with the reason. A report that shows only the hits is a sales document and gets read as one.

WHAT THIS CHANGED ABOUT THE PRODUCT

The report became narrower and much harder to argue with

We claim less than we did a year ago. We no longer say a client is invisible, because that is demonstrably untrue and the client checks it. We show which gap sits in their own pages, who fills it, and what comes out. That is less spectacular and it survives a boardroom.

// who does what, and what each one catches

The machine finds the volume. A person decides what it means. Both are wrong regularly.

Every finding in this document was produced by a person and a model working the same case. That is not a staffing note, it is the product. We sell the claim that machines read confidently and get facts wrong. We cannot deliver that claim with a process that is only a machine.

The machine does the volume

Everything that has to be read twice and compared

Every page fetched with and without JavaScript, every sitemap read to the end, every business fact diffed against every other instance of itself on the domain, and every cited document opened to see whether its text layer holds what the citation claims. On the most recent scan that was seven pages, two sitemaps, two outcome reports and twelve agent answers in one evening.

A person makes three decisions

The three things no loop produces

Which object carries the revenue. Which fact is worth chasing. And whether a citation actually holds. Those three decide what the report is about, and they are what it is bought for. The first one alone has moved a scan from admission tickets to a €299 suite, and from a beer to a claim printed on the bottle.

What each caught the other doing

Both directions, on one scan, in one evening

An agent produced a placement rate in quotation marks with the exact URL of a real file on the client's own domain. The number is not in the document. A person fetched the file and checked, and that is now a finding in the client's report. In the other direction, a cover page overflowed its A4 box by three quarters of a pixel; the measurement caught it before the document went out.

The number we publish

Eighteen predictions written down in advance. Nine of them wrong.

Across the three most recent scans we recorded what we expected before asking a single agent, and we were wrong about half the time. Both parties, repeatedly. None of it reached a client, because the discipline is not that one side is right, it is that each side has to answer to the other before anything ships.

The question this answers. Why can a client not point a model at their own site and ask it the same thing? They can, and it will find real problems. It will also hand them a confidently sourced figure that does not exist, and nothing in the answer will indicate which is which. Telling those two apart is the work.

// the first four, four kinds of damage

What a finding looks like

Seven completed scans. These are the first four, four different business types, and for each the finding that touched the money. All are anonymised, and all were found by placing two things side by side that the client had published themselves.

PRIVATE BANK • 77 YEARS

The page title and the company's own timeline contradict each other

The title and the og:title say "since 1804". The meta description, the first paragraph of the homepage and the company's own timeline say 1881. 1804 belongs to the family's textile business, not to the bank. A machine that carries the title away makes the bank 77 years older.

found at step 3 • two versions of one fact on one domain

MUSEUM • 5 PRICES FOR ONE ROOM

One word, "from", and nine answers that filled in the rest themselves

The most expensive room at this museum costs €299 on weekdays and €325 at the weekend. Its own page says only "from €299". Nine answers from four agents gave €235, €264.50, €277.50, €299 and "between €150 and €350". Not one gave the weekend rate.

found at step 4 • an escape word inside a price

EXPERT NETWORK • 0 OF 16

The only client case was in Latin, so agents invented one

The only case on the domain was placeholder text, and the five client logos sat there as images with no text. Across sixteen answers, not one client was ever offered as evidence. Every agent produced something different when asked what had actually been delivered.

found at steps 3 and 4 • placeholder text plus invisible proof

REGULATED MANAGER • 5 OF 5

The one condition that was not published, everybody supplied

Everything this manager publishes travelled accurately. What was not there was who may invest and from what amount. All five agents supplied a €100,000 threshold anyway, and one quoted a statutory warning in the client's name that appears nowhere on their site.

found at step 2 • a fact the buyer needs that was published nowhere

// three more scans, at brand scale

The same seven steps against companies with a compliance department

Three scans run to prove the method does not depend on the client being small. All three are anonymised, all three were found the same way, and one of them came out largely in the client's favour, which is reported here in the same detail as the other two.

GLOBAL BREWER · 1 CLAIM, 4 FIGURES

The alcohol content of one product, published four ways on the company's own hosts

The same drink is declared at three different strengths across the brand's own market sites, and a fourth figure appears on a fulfilment host. A fourth agent reached across to a page written for another market to tell a reader in a medically sensitive situation that the product was completely safe, with every sourcing step performed correctly.

found at step 3 · one fact, four values, none marked as the current one

RETAIL BANK · 11 OF 12

The emergency number held. The page a customer starts from never mentions it.

Asked for the number to call when a card has been stolen, the agents returned the correct emergency line eleven times out of twelve — the best result we have measured on a question that matters. The finding sits next to it: the customer service page does not name that line at all.

found at steps 3 and 5 · the well built page wins, and its neighbours lose to directories

CODING BOOTCAMP · 3 PAGES, 4 OFFERS

One programme, live at three addresses, on different commercial terms

The same course is published at three URLs, all three submitted to crawlers in the client's own sitemap, with two prices, three durations and two different money-back windows between them. A second agent supplied a placement rate in quotation marks with the exact URL of the client's own outcomes report, where that figure does not appear.

found at step 3 · old versions left live and in the sitemap are not an archive, they are a second offer

// four promises

What this project refuses to make prettier

Measured

No recommendation without evidence

Every finding traces back to a fetch we can repeat, or to a line in the client's own log. What we cannot observe, we call unverifiable. Agent answers are recorded and shown, but never carry a finding.

Falsifiable

We name the layers that are theatre

An emerging standard with no measured consumer is reported as exactly that, even when a green tick would close the deal faster. We build llms.txt for free, report it as unread, and never invoice it.

No cloaking

Nothing may make the site worse for people

No content for agents only, no hidden text, no second version of a page. Nothing we build may contain a fact that is not in the rendered page a human sees. Enforced in code, not in a policy document.

Refund

If a finding cannot be reproduced, the report is refunded

Everyone in this market sells a promise about the future. This is a measurement of the present, which is why we can give that guarantee. It is on the last page of every report.

// who it is for

From sole trader to listed company, weighted to the European middle

Dutch adoption sits at 22.7% of companies with ten or more staff, but at 17.8% at the small end against 59.2% at five hundred plus, and the main barrier is lack of experience at 74.6%, not cost at 22.8%. Only 7% of Dutch SMEs consider themselves well informed about the AI Act, while 30% already use generative AI.

The seven completed scans covered a museum, an expert network, a regulated asset manager, a private bank, a global brewer, a retail bank and a coding bootcamp. Seven business types, seven different revenue objects, and in all seven cases the same seven steps.

// the offer

Three tiers, and the first is free

Nobody should have to pay to discover that they have a problem. The scan is the product the conversation starts with, and it is complete, dated and the recipient's, even if they buy nothing else.

HOW AN ENGAGEMENT RUNS · EXCLUDING VAT	
1 · The scan — the diagnosis, proven from the client's own pages	FREE · € 0
2 · Implementation report — every URL, code and copy ready to paste	CREDITED · € 750–5,000
3 · The blocks — we do the work, verification re-measurement included	FIXED · € 750–44,500

Fixed price, no hours

Difficulty sets the tier, not duration

Every block is Quick, Medium or Extensive: how much judgement it takes to get right. There is no hourly rate, no day rate and no lead time. Every block is quoted before it starts and does not move afterwards.

What you get

One document, and it is complete

Fourteen to seventeen pages, in Dutch or English. It opens with why this costs money, then the tasks a real visitor came to do with a verdict per task, then the findings with the evidence beside them, and it closes with a price and three things the client can do for free themselves.

Blocks with a name

Named after the risk they remove

Disclosure Alignment · Machine Readability · Business Record · Rights & Permissions · Content Currency · Claim Currency · Proof Visibility · Retrieval Access · Price Delivery · Estate Scan · Verification re-measurement.

One thing we deliberately do differently from a year ago: we now do ask agents what they say about a client, and those answers appear in the report. Never as a finding, always as illustration, always with the follow-up question about the source, and always with our own model in the same table.